

Source of the Dataset

- <https://www.kaggle.com/uciml/breast-cancer-wisconsin-data>

```
In [1]: import pandas as pd
df = pd.read_csv('cancer.csv')
print(df.shape)
df.head()
```

(569, 33)

Out[1]:										
	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	points
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	

5 rows × 33 columns

```
In [2]: df.dropna(axis = 1, inplace = True)
df.shape
```

 $(569, 32)$

```
In [3]: df.head()
```

Out[3]:										
	id	diagnosis	radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	points
0	842302	M	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	
1	842517	M	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	
2	84300903	M	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	
3	84348301	M	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	
4	84358402	M	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	

5 rows × 32 columns

```
In [4]: # Separate Independent (X) and Dependent(y) features

y = df['diagnosis']
X = df.drop(columns = ['id', 'diagnosis'])
```

```
In [5]: X.head()
```

radius_mean	texture_mean	perimeter_mean	area_mean	smoothness_mean	compactness_mean	concavity_mean	concave points_mean	symmetry_mean
0	17.99	10.38	122.80	1001.0	0.11840	0.27760	0.3001	0.14710
1	20.57	17.77	132.90	1326.0	0.08474	0.07864	0.0869	0.07017
2	19.69	21.25	130.00	1203.0	0.10960	0.15990	0.1974	0.12790
3	11.42	20.38	77.58	386.1	0.14250	0.28390	0.2414	0.10520
4	20.29	14.34	135.10	1297.0	0.10030	0.13280	0.1980	0.10430

5 rows \times 30 columns

```
In [6]: # check if the dataset is balanced or imbalanced
        y.value_counts()
```

```
Out[6]: B      357
        M      212
        Name: diagnosis, dtype: int64
```

HoldOut Validation Approach - Train-test split

```

In [7]: from sklearn.model_selection import train_test_split
        from sklearn.tree import DecisionTreeClassifier
        from sklearn.metrics import accuracy_score

        X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.30, random_state = 0)
        clf = DecisionTreeClassifier()
        clf.fit(X_train, y_train)
        y_pred = clf.predict(X_test)
        print("Accuracy ", accuracy_score(y_pred, y_test))

```

Accuracy	0.9239766081871345
----------	--------------------

K Fold Cross Validation

```
In [8]: from sklearn.model_selection import KFold, cross_val_score
import numpy as np
kfold_validation = KFold(n_splits = 10)
clf = DecisionTreeClassifier()

results = cross_val_score(clf, X, y, cv = kfold_validation)
print(results)
print()
print("Results = ", np.mean(results), "+/-", np.std(results))
```

```
[0.9122807 0.9122807 0.87719298 0.96491228 0.9122807 0.96491228
 0.89473684 0.98245614 0.92982456 0.89285714]
```

Results = 0.9243734335839597 +/- 0.03350584048302555

Stratified K Fold Cross Validation

```
in [9]: # For imbalance dataset
from sklearn.model_selection import StratifiedKFold
skfold = StratifiedKFold(n_splits=5)
clf = DecisionTreeClassifier()
results = cross_val_score(clf, X, y, cv = skfold)
print(results)
print()
print("Results = ", np.mean(results), "+/-", np.std(results))
```

```
[0.90350877 0.90350877 0.92105263 0.94736842 0.90265487]
```

Results = 0.915618692749573 +/- 0.017314352857972278

Leave One Out Cross Validation (LOOCV)

```
[10]: # not recommended for large dataset
from sklearn.model_selection import LeaveOneOut
clf = DecisionTreeClassifier()
leave = LeaveOneOut()
results = cross_val_score(clf, X, y, cv = leave)
print(results)
print()
print("Results = ", np.mean(results), "+/-", np.std(results))
print("Length of Results = ", len(results))
```

[illegible]

```
Results = 0.9226713532513181 +/- 0.2671121995355125
Length of Results = 569
```

Shuffle Split/ Repeated Random train-test splits

```
[11]: from sklearn.model_selection import ShuffleSplit
ssplit = ShuffleSplit(n_splits = 10, test_size = 0.30)
clf = DecisionTreeClassifier()
results = cross_val_score(clf, X, y, cv = ssplit)
print(results)
print()
print("Results = ", np.mean(results), "+/-", np.std(results))
```

```
[0.92397661 0.94736842 0.88304094 0.94152047 0.92397661 0.94736842
 0.92982456 0.9005848 0.9005848 0.92397661]
```

Results = 0.9222222222222222 +/- 0.02043439205807736

```
In [ ]:
```